

A course in language creation



# *Conlangs University*

## Documentation 1

Sketch grammar & romanisation



April 2020

Slorany, Tutor

Romanisation chapter by Astianthus

This entire document will be working from the assumption that you are working on a language situated within a fictional world. This world needs not be documented to any extent, only to exist for either you, or your speakers. It may be set on an alternative timeline, an alternative Earth, or a completely invented world. Some parts will necessarily not apply to languages that are either intended as auxiliary languages for our world, or exist in a complete vacuum.

# Introduction

Documenting your language has two main purposes: displaying the language for others, so that you may get feedback, and keeping track of it for yourself, so you have a convenient document to refer to when *using* the language, for instance while translating or producing original text.

You must thus ask yourself the following: do you intend to teach this language to others? To yourself? Is it only for self-reference?

Then come the secondary purposes of documenting your language, in the form of several considerations that must be balanced as you see fit: usefulness/ease of use, understandability, and cultural relevance.

There are also challenges that grammars of conlangs do not commonly face, such as accountability to the speakers or to the academic community; and there are new challenges that arise, as a conlang has no native, or even fluent, speakers who could help analyse the language and provide a corpus.

As not all the sources and references for this document are freely available online, here are those I cannot provide a direct link to:

- [Catching Language](#), Multiple authors, Mouton de Gruyter (2006)
  - Specifically, chapters 1, 3, 6
- [Essentials of Language Documentation](#), Multiple authors, Mouton de Gruyter (2006)
  - Specifically, chapters 1, 9, 12

And two others that served as a reference and are available, but I couldn't properly fit in the document:

- [Documentary and descriptive linguistics](#), Nikolaus Himmelman (1998)
- [Seven dimensions of portability for language documentation and description](#), Steven Bird & Gary Simons (2003)

That second paper gives advice about the format of grammars and general information about their social, political and technological contexts, contents and value.

## What/how to document?

Conlangs are obviously quite different from natural languages in that there is no real accountability in the process: you do not have to provide evidence of your claims about how the language operates because you (whether you as an individual or as a group of co-creators) are the sole possible source for the language. of course, this gets blurred with collaborative pidgin projects such as [Viossa](#) (also [on FrathWiki](#)) or collaborative emergence projects (where all communications are spontaneous and reusable from one speaker to another as a consensus on the meaning is reached without communicating it in English) such as Lxelxe, where the documentation will necessarily rely on an input any one participant cannot fully control.

There is a lot that is possible to document for a conlang. As you are creating the culture of the speakers of the language, you may want to keep it all in the same document so that anyone who reads it has a sense of the whole project.

# Why document?

There are many reasons to document a natural language. But what about a conlang? Many posit that conlangs are useless and that efforts would be better spent documenting existing languages that are fading out of use and existence. This is missing the point that conlangs are only rarely scientific or academic endeavours, but more often artistic ones. A conlanger doesn't necessarily have the skills to properly document a natural language while maintaining accountability to its speakers and to the academic community.

This answers the primary question of "why even conlang?".

Now, let us address the question of "why document your conlang?".

Some reasons are obvious and directly useful to the creator(s): no one has a perfect memory, and writing things down for later reference is useful. It ensures that later efforts on the same project will not be redundant, or rendered less fruitful because one forgot to take into account some part of the language. Documenting maximises productivity, and minimises errors in reference.

It's also satisfying to see the documentation grow as you work, and as the amount of information (and your knowledge!) about the language increase(s).

If your conlang is to be learned by anyone, be it actors on a TV show, players of a roleplaying game, readers of a book, it is much better to have some reference of it, for both teacher and learner.

And if you intend to make your work public at any point, presenting it in a consistent and cohesive format is better than hastily fumbling to get your notes together.

Another aspect of it, often neglected, is inspiration. A clean document is motivating, both for the author(s) and for anyone reading it after the fact. Seeing people publish well-written, well-formatted, well-organised documentations of their works is inspiring! Who doesn't want to produce beautiful results?

## Goals (again)

It's important to know what you are aiming for when documenting your language. How do you want to present your language? It can be a few things:

- A short document outlining only the essentials of the grammar;
- An extensive grammar detailing as many of the intricacies as possible;
- A phrasebook-style guide for learning the language, similar to those for tourists visiting a country;
- Whatever else you can think of.

For the sake of brevity, this course will focus mainly on laying out the contents for a *sketch grammar*, which corresponds to the first item on the list above.

### **A note about copying another grammar's outline:**

It is tempting to just take the outline of a natural language's grammar and filling the parts, chapters and sections with information about your language, however no two languages are the same.

Obviously, there are caveats to such an approach: no two languages are similar enough to warrant documenting their features in the same order, unless you are knowingly limiting accuracy and are aiming more for a comparative analysis of, for instance, two dialects of the same language.

My recommendation is to simply not look at another language's grammar for guidance, at least not for the overall organisation. Looking at them is of course a good idea if you are struggling with page layout.

# The considerations

This chapter is a quick look at the considerations mentioned in the introduction of this document.

## Audience

Regardless of who the intended audience for your documentation is, it should be comprehensive and understandable to *you*, the author.

However, you have to keep in mind that what *you* know about your language — because it is your creation and occurred in your own mind — when reading any part of your documentation is not necessarily what another reader will know: you will have to make sure your readers are aware or can be aware, at any time, of everything that is used in any example inside of your documentation, as well as how every part functions, at least on a basic level.

A language documentation will be written differently depending on who is intended to read it: are they language specialists (e.g. linguists)? Are they language teachers with no linguistic training? Are they (potential) learners of the language?

Who you intend your audience to be *should* shape how you write your documentation. A documentation by a linguist for another linguist will necessarily be different from one that is written by a linguist for laymen, and in turn from one by a layman for other laymen.

Who your audience is will also be determined by the *point of view* you will choose to write from, a point discussed further in this document.

## Ease of use

Whoever your audience is, the sketch grammar should meet their needs and enable them to get an overview of the main traits of the language. The writing should be clear and to the point.

All terms that may not be known to your audience (i.e. whose usage is mostly confined to linguists and other language specialists) should be explained, as should all terms that have variable definitions (adverb, particle, adjective, noun...).

## Understandability

The documentation for a language, especially that of a conlang as it is inherently non-academic in nature, should ideally find a balance between providing an accurate understanding of the language, and providing knowledge of the language to as many people as possible.

It is very likely that the format we will be exploring most here will exclude anyone with no training in linguistics, self-taught or not, by necessity.

How much information you should include has a floor, but no ceiling: you will *need* to include everything necessary to understand all the utterances in the documentation, but you can always choose to create more utterances to document.

You tell us what our documentation **should** have, but is it really mandatory?

No! Nothing here is *mandatory*: they are suggestions about what I think is best. This is based on the assumption that you are not a trained linguist and are not seeking to fully emulate fieldwork or the analysis of a corpus, as creating enough content for this would likely take you years.

Please do feel free to ignore any and all recommendations that do not fit the vision you have for your own creation.

## Universality

It can be hard to make a difference between what is needed and what isn't, as it can vary from one utterance to another.

For instance, for a sentence meaning "she gives him a kiss on the cheek", it is generally useless to say that a human has lips, that a kiss is applying those lips to something, that humans have cheeks, that a kiss is affectionate...

Those small things are perfectly obvious, at least in anglophone cultures (and as English is the language I am using here, I am assuming you are familiar with that), and can be omitted.

## Terminology

### Problem 1

Every language is unique and, as such, should be treated uniquely and specifically.

### Problem 2

All languages have elements in common with other languages, elements that behave similarly and to varying degrees of difference.

Thomas Payne's *Towards a balanced grammatical description* gives a few good examples of these problems in its chapter 3 *Inclusiveness vs Usability*.

---

\*

---

In order to give a reader of your documentation an approximate idea of what to expect from a given element of your language, it makes sense to give it a name that is already used for a similar feature in another language known to the reader.

For instance, "noun" or "verb" as parts of speech should be familiar to most people reading and, despite their not having a single definition cross-linguistically, they are convenient approximations.

For these types of very broad concepts, getting "close enough" is a valid approach as long as it is followed by an explicit definition of what exactly a label with a broad name entails.

---

GHO

---

A balance can be found by only coining specific or new labels for the things in your language that are more unique than others, highlighting them this way.

Equally, it is good to first define something, before assigning it a label: the description you come up with may be quite different from the one of a label you had in mind for it at first, or you might be driven away from your original intents by wanting to stick to the label's definition in English (or another language you know).

This is a good reason why established linguistics are sometimes detrimental to the conlanging process, in that they can limit it.

## Cultural relevance

As said, in the section above, there is no limit to how much information you can put into your sketch grammar.

However, as this is a conlang with no *existing* people who speak it, there is no reference available for these people and you will have to create them.

Generally, grammars (and sketch grammars) have a section in their introduction that gives information about the language as a whole: where it's spoken, how healthy it is, other names for it, and who speaks it.

But this is not a natural language with speakers and a culture you can inform yourself about through the use of an encyclopedia or an internet search: this is a constructed language, and you may want to explain how the language relates to the culture of its speakers.

To do so, you will have to either include it in your introduction, or in a specific part of the documentation, or in a separate document.

If you are an avid worldbuilder, you may already have that separate document.

## Point of view

Depending on how your conlanging relates to your worldbuilding (and *if* it does relate to some), there are several points of view you can adopt for describing the language.

Where you choose to describe the language from will be a factor in how you choose to describe it.

Conlangs are peculiar to describe as there is no way to tell if the documentation for it is describing *all* of the possibilities of the language, or even how much of it: it probably isn't, but we do not have any other speaker but its creator(s) and author(s) who could verify or deny that.

### Inside & outside

Here, you would be writing the documentation as someone in your world, unaware of ours.

You may choose to adopt the point of view of a character in your world, who speaks the language natively.

- In this case, who is this speaker trying to inform? Why?
- Are they trying to document their own language for reference by other native speakers?
- Are they trying to communicate it, and its intricacies, to language specialists from another culture?
- Are they trying to teach it?
- Are they writing a *description* of it, or a prescriptive manual, dictating usage?



Another possibility is to have your documentation be written from the perspective of an outsider, a native speaker from another language.

You can ask yourself the same questions as with the native speaker as author, but consider the following as well:

- How (well-)informed are they about the language they are trying to document?
- What are their biases about it?
- How different is their own native language from the one they are writing about?

## From the top

You could also be documenting this language as your own, whilst acknowledging that *you* are not a native speaker of it, and describing how the fictional speakers use it.

In this case, you can choose to deliver information as though it were absolute, acting as an omniscient entity. This will mean that all the utterances you will list can be considered as the only valid ones, until you provide more.

---

---

\*  
\*\*

---

---

In the end, where you choose to describe the language from is only a matter of how much of the worldbuilding you want to give away, and how.

It is of course possible to mix and match the different approaches at your leisure, or to imagine others. This is in no way exhaustive, complete, or even close to perfect.

# Sketching it out

## What is a sketch grammar?

In academia, a sketch grammar can be a few different things:

- a preliminary work that will be the base for a larger description;
- a summary of a larger description;
- an introduction to the language's grammar's essentials at the front of a research paper;
- the grammar section in a dictionary.

Note that those four are not identical and can be built in different ways. For instance, a sketch grammar is more likely to have a small lexicon attached that only contains the words necessary to decipher the examples plus some more that are commonly used and/or essential to basic expression in the language, while a dictionary grammar will use that grammar section to enable some economy: if a productive affix with a regular behaviour exists in the language, it will be easier and more economic to describe it once in the grammar than to list in the dictionary all the words it can exist in.

Grammars at the front of research about a specific topic will likely make without any aspects that are not necessary to the understanding of the topic and of the examples in the paper.

Sketch grammars of conlangs are typically modelled after the first type listed above: a preliminary work that will be the base for a larger description.

This is partly by necessity: you will likely be building the documentation as you go along since this is entirely a creative endeavour and, at the time you will be writing this sketch grammar, might not be based on any substantial content that would enable you to conduct an analysis of a corpus, instead of creating the corpus for the grammar.

A sketch grammar is thus not a "short grammar" as it does not aim at describing the language as a system but at facilitating the access of the documentation.

---

\*\*\*

---

This last sentence above leaves a rather large problem for a conlang's sketch grammar to deal with: how do we make it seem like this language *exists*, when it doesn't? It has no speakers and exists only in the mind of its author.

A solution (and not the only one, but the one I choose to focus on in this document) to this problem is to use a lot of examples.

---

\*\*\*

---

Let us then proceed with the assumptions that:

- you are indeed documenting the language as you create it;
- you do not have a large corpus (yet);
- you have defined the goals of your language (see the [Introduction course of this series](#), and *The Greatest Conlang Ever*, also in this series);
- you have defined the goals for your documentation (see earlier in this document).

## Examples & glossing

Examples will constitute the basis of your conlang's sketch grammar, and it will need to contain a good number of them.

When using [interlinear glosses](#) (generally following the [Leipzig Glossing Rules](#)), you can shove any amount of information into the abbreviations in the glosses, from minimal to maximal.

Let us take an example used in *Essentials of Language Documentation*, Chapter 12:

```
bona moon  
ART woman
```

Here, **bona** is simply indicated as being an article. As such, there is an expectation that the entire paradigm of articles will be explained in the documentation, detailing the information carried by that specific article: Does it express noun class/gender? Number? Definiteness?

```
bona          moon  
ART:A_CLASS.SPEC.SG.OBJ woman:A_CLASS
```

Doing it this way means the sketch grammar can do with only explaining the various abbreviations. Here, a specific singular object article agreeing with a noun that belongs to the A class.

The first option is generally preferred as it avoids overly long glosses, which can cause even simple sentences to take up several lines to type out with their gloss.

Whichever option you choose to use, you should do so consistently.

## Contents

What a grammar *should* contain and how it should be written have been the topics of a few critiques and standardisation proposals, such as the [1977 Lingua Descriptive Studies Questionnaire by Bernard Comrie & Norval Smith](#), [Thomas Payne's \*Toward a balanced grammatical description\*](#) or — perhaps even more famously in conlanging circles — [Van Driem's \*Dallas Manifesto\*](#).

However there are a few basics most authors, if not all of them, agree should be in a grammar of any type as they form the basis for an understanding of a language for linguists and language learners alike:

- the basics of the sound system and orthography;
- how simple sentences are formed;
- what major word classes there are and how they are distinguished;
- how word groups (noun phrase, verb phrase, adpositional phrase...) function;
- inflectional morphology;
- derivational morphology.

These are sufficient for a sketch grammar which, since there is no massive corpus for a conlang and you have to create every utterance you will document, is the most likely format for a documentation of your creation. Note that putting them in the order above is not necessarily the best: you will have to devise an order that is most adequate for your language.

Together with a dictionary, the sketch grammar should be sufficient to understand and loosely translate a simple text in the language.

**NOTE:**

The Lingua Descriptive Studies Questionnaire is **not** to be considered as a checklist, and neither are any other questionnaire or other resource, beyond the absolute basics (for instance, the 5 main sections of the LinguaQ).

Consider them as vague directions of "stuff languages may do", not absolutes or imperatives.

Perhaps the best set of guidelines for a sketch grammar is [the one from Mark Donohue](#): being short, it only has to address the essentials and as such does not provide unnecessary constraints.

I won't be reiterating too much of the advice given in it, so I greatly advise you read it. Instead, the following sections will address problematics that are either absent from it or too shallowly described for conlanging.



The various articles cited and linked throughout this lesson bring better descriptions of the possibilities for organising the contents of your grammar than I would be able to, lest I simply regurgitate their contents.

I would nonetheless advise that your introduction, or some form of preface if you elect to not write the grammar from your very own perspective, be different from that of a natural language and evoke your goals, intents, and methods, for both the language and the world around it, if applicable.

# Tools, ways, & methods

Here, I describe strategies and tools I have employed or currently employ for my own languages and creative processes.

These are things that have worked for me and, while I will be presenting quite a few, I cannot guarantee they will work for you.

For everything during the ideas and research notes stage leading up to the first draft, I use paper exclusively. Lots of paper, from writing on post-its to printing.

After this, I switch to LaTeX, described further down!

## Storing words

When I craft words, I store them in a spreadsheet for quick reference. One tab of the sheet will have some research capabilities built into it, such as a space to let me type any word in English or in the conlang, as a way to check whether I already have something for a given concept or already make use of a given word.

Another tab will have all the words and information about them, organised as such:

| ↓<br>Words | Grammatical<br>category | Other<br>info 1 | Other<br>info 2 | Definition<br>1 | Definition<br>2 | Definition<br>3 | Definition<br>4 | Notes | Used<br>in... |
|------------|-------------------------|-----------------|-----------------|-----------------|-----------------|-----------------|-----------------|-------|---------------|
| A1         | B1                      | C1              | C1              | C1              | C1              | C1              | C1              | C1    | C1            |
| A2         | B2                      | C2              | C2              | C2              | C2              | C2              | C2              | C2    | C2            |
| A3         | B3                      | C3              | C3              | C3              | C3              | C3              | C3              | C3    | C3            |

This allows me to store the words in a way that's easily searchable.

**Other info 1 & 2** are about absolutely anything specific to these words. Are they restricted in usage? Are they often "misused" by youths in a way that annoys their grandparents?

**Defintions** are not just English equivalents or translations: here, I actually define the words, talk about what they mean rather than how to translate them.

**Notes** may seem like a redundancy when "Other info" exists (twice!), but it isn't: here I put notes for myself as a creator: did I intend for this word to be used often in translations? Does it have synonyms? Antonyms?

**Used in...** is the place where I type all the sentences and examples I've used the word in.

Is this ideal? No. Will this work for you? No guarantee.

But it works for me, and it can transition into the next step pretty well, with the help of some automation.

Here is [a link to a template in Google Sheets](#), if you wish to try it out.

Some people [will tell you that the database format sucks for storing words](#), and they're perfectly right! I, much like the author, despise that system when it comes to translating content or simply looking at it. It's terrible, but the sole purpose of that system is to be efficient, not to be good.

Plenty of applications can be used to store words, and even some part of your grammar (in **bold**):

- [Lexiconga](#)
- [Polyglot](#)
- [SIL Fieldworks](#) and other [SIL software](#)
- [Lexique Pro](#)
- [WordTheme](#) (Android)

These applications are convenient and allow you to always have your material available on the go, but in the end nothing beats an actual dictionary format. See below!

## Formatting

Here, we are talking about a final product, something you can publish and display.

While I could spend a lot of time going on and on about how and why LaTeX is the best thing ever, I'll let someone else do it for me:

- [Part I](#)
- [Part II](#)
- [Part III](#)

In short, LaTeX is great because it's ***relatively*** easy to learn and produces amazing and consistent results. There exist free services so you don't even have to install anything to get started, as everything can be done online. [Overleaf](#) is the most popular, and it offers a decent [tutorial](#) as well as [quite a few templates](#) organised by category.

Here is a [PDF document made with LaTeX](#), as an example, from my own creative writing.

Here, "***relatively*** easy" means that, while there is a learning curve, it is not particularly hard to get good-looking basic documents, especially through the use of templates.

However, for anything involving some amount of customisation, LaTeX can be very hard to make work *the way you want it to*.

The inability to see the changes as you make them can also be a big impediment.

If you prefer the convenience of a [WYSIWYG](#) editor, use that. **LaTeX is good, but it might not be worth the time you will spend learning it!**

But I would be remiss not to mention that regular text editors can also be used!

LaTeX, while being relatively easy to learn, has several drawbacks:

- you still have to learn it, no matter how easy it may be, it's going to take some time;
- other text editors, while not *as* capable as LaTeX, are *capable enough* for the purpose of creating conlangs and easier to handle;
- you are probably already familiar with the more common text editors.

Even David J. Peterson, creator of Dothraki, uses Pages on Mac, and not LaTeX.

If you fancy something with more organisational tools built in, you can use Scrivener, a popular (commercial) programme for writing novels or screenplays, or [yWriter](#), which is free.

## Writing the language

While this document will not cover how to create a native writing system for your language and its speakers, it needs to discuss how to represent the language's words without having to type everything with IPA symbols: a romanisation.

# Romanisation

This part of the course was written by Astianthus.

It would be a pain to use the IPA all the time when writing in or about your language. For many languages, such as English, Hindi, and Korean to name three, this is a non-issue, as they have an **orthography**, i.e. a canonical written form. Some languages, however, aren't used in writing by their native speakers and don't have an orthography (this is likely the case for your conlang right now). In these cases, we get to decide how we write the language down! If we decide to use some variation on the Latin alphabet for this (like English and the majority of European languages do in their orthographies), this written representation gets called a **romanisation** (or romanization).

The term 'romanisation' is also used for Latin alphabet-based systems for writing languages that have an orthography in another script, such as writing Mandarin Chinese with [pinyin](#). Such a romanisation should not be confused with a **transliteration** (to the Latin alphabet), which is when another writing system is represented with the Latin alphabet in such a way that there is a more or less one-to-one relationship between the scripts' letters. This topic will be about creating a romanisation for an unwritten language, but most points are also applicable to creating a romanisation or a transliteration for a written language.

Because we're free to adapt the romanisation to our purposes, we can avoid a lot of the peculiarities of some real-world orthographies such as silent letters, ambiguous pronunciation, and so on, which are mostly due to historical reasons and language change. A conlang can of course have a fictional orthography in addition to or instead of a romanisation, which may very well mimic the complexities mentioned, but this will be covered in another topic.

## Good romanisations

A romanisation should be easy to use (that's why it exists), but if you plan to show your language to others it should be easy to read and understand as well. Some general guidelines for doing this, arranged roughly in decreasing order of importance, are the following: the romanisation for your conlang should ...

- Align with your own preferences and goals (you'll probably use your romanisation a lot, so it's important that you like it yourself).
- Represent the phonetics of your language straightforwardly and unambiguously. Specifically:
  - Each letter or letter combination should always represent the same phoneme.
  - Each phoneme should always be represented with the same letter or letter combination.
  - Silent letters should be avoided.
- Use letters and letter combinations for their usual sound values as much as possible.
- Use its diacritics consistently, if it has any (i.e. avoid having one diacritic for multiple different purposes).
- Be easy to type.

One tendency among some conlangers is to make a very exotic-looking romanization (often using gratuitous apostrophes and diacritics) for no reason other than to look alien. This can of course be effective if you want your conlang to look a certain way to outsiders, but it's generally frowned upon by the conlanging community.



Consider making a constructed writing system for your language if you want to give it a certain visual aesthetic.

The next sections will elaborate on some of the points in the list above, mainly recommendations for how you can pick what symbols to use for your specific sounds.

## What symbol should I use... for this consonant?

The main part of designing a romanisation is of course choosing what symbols to use. In general, using the IPA as a starting point is a good idea. Wikipedia is a good resource for checking how a sound is written in lots of languages ([example](#) for ɬ). To create a reasonable orthography, no more tools or knowledge is really needed, but there are a few patterns and solutions to common problems that I feel are worth stating explicitly. This section is focused on consonants, and the next is focused on vowels. Some notes on specific consonants or consonant groups follow in no particular order.

The velar nasal [ŋ] is often represented with the digraph *ng* as in English, or by *ŋ* as in the IPA, and the languages [Wolof](#) and [Northern Sami](#). Other options include *g*, if that's not being used for anything else, or an accented version of *n* such as *ñ*, *ń* or *ň*. I personally dislike using an accented version of *n* as the ones mentioned usually represent [ɲ].

Following the latin orthographies of many Slavic languages (such as [Gaj's Latin alphabet](#) for Serbo-Croatian), the postalveolar fricatives and affricates [ʃ ʒ tʃ dʒ] are often written *š, ž, č, and dž*. Other practices are sometimes adopted from the same source, such as writing [ts] as *c*, and [ɟ ʒ tɕ dʒ] as *ś, ź, ć, and dź* when they contrast with postalveolars. Taking inspiration from English and writing [ʃ ʒ tʃ dʒ] along the lines of *sh, zh, ch*, and *j* isn't uncommon either.

Fricatives that aren't easily romanised in another way are often written with a digraph of a corresponding plosive + *h*, so [ɸ β θ ð x ɣ] are often seen as *ph, bh, th, dh, kh, and gh*. Of course, if there is no need to distinguish [ɸ β] from [f v], they can simply be written as *f* and *v*. Another use of *h* in digraphs is for voiceless versions of sonorants, so [ɱ ɲ ɭ ʀ] are often written as *mh, nh, lh, and rh*, or the other way around as *hm, hn, hl, and hr*.

Pinyin writes the voiceless unaspirated [p t k] as *b, d, g*, reserving the usual voiceless symbols *p, t, k* for the aspirated [p<sup>h</sup> t<sup>h</sup> k<sup>h</sup>] (with a corresponding system for the affricates). This idea is worth copying if your language, like Mandarin, distinguishes plosives by aspiration and not voicing.

If there are two r-like sounds and one is a trill, a natural way to disambiguate these is by writing the trill with *rr* and the other with *r*, as in [Spanish](#).

An apostrophe is very commonly used to represent the glottal stop [ʔ] or glottalisation of any kind. In particular, ejectives are often written as *t', k'*, and so on, mirroring the IPA. Another common way of romanising the glottal stop is with any left-over letter such as *h, x*, or commonly *q*. Some choose to use the actual symbol ʔ.

The same general principles apply to vowels as well as consonants, but there are some additional things to keep in mind when it comes to vowels, which mostly stem from the following facts:

- There are fewer obvious vowel letters, so diacritics and unconventional vowel letters are often used to get more options.
- Vowel letters are often marked for categories such as length, stress, tone, and nasalisation. This usually means more diacritics.

Among conlangers, the usual vowel letters *a*, *e*, *i*, *o* and *u* are in general expected to correspond roughly to the IPA vowels [a e i o u] rather than their English pronunciations, which are only rarely the same as the IPA. When vowels other than those mentioned are needed, there are three common solutions:

- Using a letter with a diacritic. A common example is using *ü* and *ö* for [y] and [ø].
- Using an unconventional letter, such as using *ı* for [u] as in [Turkish](#). Simply copying the IPA character for the vowel would also fall into this category. For example, [ɛ ɔ] are often written as *ε* and *ο* in African languages that use the Latin alphabet.
- Borrowing a letter usually used for consonants, such as using *w* for [u] as in [Welsh](#) or *v* for [ə] as in romanised [Cherokee](#).

If some of your vowel letters have to be accented, it may be worthwhile to think about which vowels are common in your language. If you have letters for both [e] and [ə], writing them as *e* and *ě* might seem reasonable, and it is, but if [e] appears less commonly than [ə], choosing the symbols *é* and *e* instead might be more to your liking if you prefer the unaccented letters to be more common than their accented counterparts.

If you use the same diacritic on multiple letters, it's a good idea to make the usage at least mostly consistent. As an example, [Hungarian](#) uses the trema (or umlaut) on back vowels to form a fronted version of that vowel (with the same rounding): *u*, *o* represent [u o], and *ü*, *ö* represent [y ø]. If there is some form of vowel harmony (as there is in Hungarian), a consistent writing system can sometimes make the process clearer than it would otherwise be. Simplifying a bit, in Hungarian *u* and *o* can alternate with *ü* and *ö*, depending on what other vowels exist in a word. Because of the orthography, this often as easy as adding or subtracting a trema.

As stated in the beginning of this section, vowel letters often carry more information than what vowel quality they represent. You can think of vowel quality as the "base" IPA symbol used for the vowel: [à] and [á:] have the same vowel quality, but differ in the categories of length and tone. The remaining paragraphs of this section on vowels will focus on how to represent such categories.

To start off, stress could reasonably be marked with most diacritics, but the most common option by far is to use an acute accent (as in Spanish).

Nasalisation is commonly marked with a tilde, as in the IPA, or with an ogonek (*ą*, *ę*, *ĩ*, *ę*, *ų*) as in [Navajo](#). A third option is to mark nasalisation with a separate letter if this does not create unwanted ambiguity: [ã] might for example be written as *an* or *am*.

Long vowels can be represented a number of ways. Perhaps the most obvious strategy is to write the vowel in question twice, which is what [Finnish](#) and Navajo do. In the case of Navajo, this is useful because vowels already carry diacritics above the vowel for tone and below the vowel for nasalisation (as mentioned above), so there's little room for a third diacritic. If adding a diacritic isn't a problem, a macron as in [Hawaiian](#), or an acute accent as in Hungarian are popular options. Hungarian also uses the special characters *ű* and *ő* as

versions of *ü* and *ö* with an acute accent. A third option is to mimic the IPA and simply write *e:* for a long *e*, and so on. However, if the readers of your conlang are unfamiliar with the IPA, this has the possibility of being quite confusing.

Lastly, there is tone, which there are two main ways to romanise:

- Using a diacritic for each tone, like pinyin and Navajo.
- Using an entirely separate letter, as is done in romanisation for [Hmong](#).

In either system, a tone may well be left unmarked if it more common than others, or the default in some other sense. Diacritics will be familiar to those who know the IPA, especially if you follow the convention that acute (*á*) is high tone, grave (*à*) is low tone, and so on. However, if you already use diacritics, stacking them runs the risk of looking very cluttered. If you don't mind comparisons with [Vietnamese](#), this could be fine. Otherwise, the second option could be worthwhile.